




PM reviewer reliability vs ground truth
(endorsed = accept+modify vs reject, scored against Point_Mass_v5)



Per-criterion agreement with GT (fraction of flags each reviewer called correctly)



How each reviewer responded to the AI's flags



PM reviewer reliability scorecard

Reviewer	session	flags	TP	FP	FN	TN	Agree	Prec	Rec	F1	modify
SME 1	cf9c5ba6	42	21	21	0	0	0.50	0.50	1.00	0.67	0
SME 2	aee204f7	40	13	18	7	2	0.38	0.42	0.65	0.51	0
SME 3	6e1e38fc	42	15	17	6	4	0.45	0.47	0.71	0.57	4
SME 4	9dc74a53	39	19	7	2	11	0.77	0.73	0.90	0.81	0